

Names of the classifiers used

SVM

Logistic Regression

Decision Tree

KNN

Naive Bayes (Gaussian)

Import Libraries

```
!pip install seaborn
!pip install pandas

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from sklearn.preprocessing import StandardScaler, OneHotEncoder, LabelEncoder, MinMaxScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.ensemble import RandomForestClassifier
from sklearn.inspection import permutation_importance
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV, StratifiedKFold
from sklearn.svm import SVC
from sklearn.decomposition import PCA
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, roc_curve, auc, precision_score, recall_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from scipy.stats import loguniform
from sklearn.metrics import ConfusionMatrixDisplay, RocCurveDisplay
```

```
Requirement already satisfied: seaborn in /usr/local/lib/python3.12/dist-packages (0.13.2)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in /usr/local/lib/python3.12/dist-packages (from seaborn) (2.0.2)
Requirement already satisfied: pandas>=1.2 in /usr/local/lib/python3.12/dist-packages (from seaborn) (2.2.2)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in /usr/local/lib/python3.12/dist-packages (from seaborn) (3.10.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.0.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.10.0)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.22.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.3.1)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (20.4)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (8.3.0)
Requirement already satisfied: pyarsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.3.1)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.7.3)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.2->seaborn) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.2->seaborn) (2025.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.7->matplotlib!=3.10.0) (1.16.0)
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages (2.2.2)
Requirement already satisfied: numpy>=1.26.0 in /usr/local/lib/python3.12/dist-packages (from pandas) (2.0.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)
```

Import Dataset and print the first 2 rows for sanity check

```
df = pd.read_excel('/content/agro_dataset.xlsx')

print(df.head(2))
```

```
Area MajorAxisLength MinorAxisLength Eccentricity ConvexArea \
0 87524        442.246011      253.291155     0.819738      90546
1 75166        406.690687      243.032436     0.801805      78789

   Extent Perimeter    Class
0 0.758651  1184.040  Kecimen
1 0.684130  1121.786  Kecimen
```

▼ Data Visualization and Exploration

```
sns.set_theme(style="whitegrid")

plt.figure(figsize=(8, 6))
ax = sns.countplot(x="Class", data=df, palette="viridis")
total = len(df)
for p in ax.patches:
    h = int(p.get_height())
    ax.text(p.get_x() + p.get_width() / 2., h + 0.01 * total, f'{h} ({h/total:.1%})', ha='center')
plt.title("Distribution of Classes")
plt.tight_layout()

numeric_cols = df.select_dtypes(include=np.number).columns
for col in numeric_cols:
    plt.figure(figsize=(10, 4))
    sns.boxplot(data=df, y=col, x="Class", palette="viridis")
    plt.title(f'Distribution of {col} by Class')
    plt.tight_layout()

miss_frac = df.isnull().mean().sort_values(ascending=False)
print("\nMissing fraction by column:")
print(miss_frac)
plt.figure(figsize=(8, 3))
sns.barplot(x=miss_frac.index, y=miss_frac.values, palette="magma")
plt.ylabel("Fraction missing")
plt.xticks(rotation=45)
plt.title("Missing data fraction per column")
plt.tight_layout()

plt.figure(figsize=(8, 6))
sns.heatmap(df[numeric_cols].corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Matrix of Numerical Features")
plt.tight_layout()

# Pairplot to visualize feature relationships by class
print("\nGenerating Pairplot...")
sns.pairplot(df, hue="Class", palette="viridis")
plt.suptitle("Pairplot of Features by Class", y=1.02) # Add a title to the pairplot
plt.tight_layout()

plt.show()
```




```
/tmp/ipython-input-541899765.py:4: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` a  
ax = sns.countplot(x="Class", data=df, palette="viridis")  
/tmp/ipython-input-541899765.py:15: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` a  
sns.boxplot(data=df, y=col, x="Class", palette="viridis")  
/tmp/ipython-input-541899765.py:15: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` a  
sns.boxplot(data=df, y=col, x="Class", palette="viridis")  
/tmp/ipython-input-541899765.py:15: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` a  
sns.boxplot(data=df, y=col, x="Class", palette="viridis")  
/tmp/ipython-input-541899765.py:15: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` a  
sns.boxplot(data=df, y=col, x="Class", palette="viridis")  
/tmp/ipython-input-541899765.py:15: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` a  
sns.boxplot(data=df, y=col, x="Class", palette="viridis")  
/tmp/ipython-input-541899765.py:15: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` a  
sns.boxplot(data=df, y=col, x="Class", palette="viridis")  
/tmp/ipython-input-541899765.py:15: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` a  
sns.barplot(x=miss_frac.index, y=miss_frac.values, palette="magma")  
  
Missing fraction by column:  
Area          0.0  
MajorAxisLength 0.0  
MinorAxisLength 0.0  
Eccentricity   0.0  
ConvexArea     0.0  
Extent         0.0  
Perimeter      0.0  
Class          0.0  
dtype: float64
```




Class


Distribution of Area by Class




Distribution of MajorAxisLength by Class




Distribution of MinorAxisLength by Class








Distribution of Eccentricity by Class



Distribution of ConvexArea by Class







▼ Data Visualization and Exploration

The correlational analysis can have an effect on feature selection. High correlation between features indicates that they are providing similar information, which can lead to multicollinearity and redundancy in the model. Identifying these relationships helps in deciding whether to remove one of the highly correlated features or use dimensionality reduction techniques like PCA.

```

correlation_matrix = df.corr(numeric_only=True)

upper = correlation_matrix.where(np.triu(np.ones(correlation_matrix.shape), k=1).astype(bool))

to_drop_highly_correlated = [column for column in upper.columns if any(upper[column] > 0.95)]

print("Features to drop due to high correlation:", to_drop_highly_correlated)

df_reduced_correlation = df.drop(columns=to_drop_highly_correlated)

print("DataFrame shape after dropping highly correlated features:", df_reduced_correlation.shape)

X = df_reduced_correlation.drop('Class', axis=1)
y = df_reduced_correlation['Class']

numeric_transformer_pipeline = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler()))]

categorical_transformer_pipeline = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])

numeric_features_reduced = X.select_dtypes(include=[np.number]).columns
categorical_features_reduced = X.select_dtypes(exclude=[np.number]).columns

preprocessor_reduced = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer_pipeline, numeric_features_reduced),
        ('cat', categorical_transformer_pipeline, categorical_features_reduced)
    ],
    remainder='passthrough'
)

X_processed = preprocessor_reduced.fit_transform(X)

print("Preprocessing complete with reduced feature set. Processed features shape:", X_processed.shape)

```

Features to drop due to high correlation: ['ConvexArea', 'Perimeter']
 DataFrame shape after dropping highly correlated features: (900, 6)
 Preprocessing complete with reduced feature set. Processed features shape: (900, 5)

Data Pre-processing and cleaning

▼ Justification of Feature Selection and Transformation Techniques

To prepare the data for modeling, we applied the following techniques:

- **Correlation-based Feature Selection:** Removed highly correlated features (correlation > 0.95) to mitigate multicollinearity and reduce redundancy, improving model stability and interpretability.
- **Principal Component Analysis (PCA):** Used in the SVM pipeline for dimensionality reduction by transforming features into a smaller set of uncorrelated components. Useful for handling correlated features and potentially improving performance and efficiency.
- **Random Forest Feature Importances & SelectKBest:** Employed these methods to identify and select features most relevant to the target variable based on impurity reduction (Random Forest) and statistical association (SelectKBest).
- **Standardization (using StandardScaler):** Scaled numerical features to zero mean and unit variance. This is crucial for distance-based and gradient-based algorithms (like SVM, Logistic Regression, KNN) to ensure features contribute equally and prevent dominance by features with larger scales.

```

X = df.drop('Class', axis=1)
y = df['Class']

label_encoder = LabelEncoder()

```

```

y_encoded = label_encoder.fit_transform(y)

numeric_features = X.select_dtypes(include=[np.number]).columns
categorical_features = X.select_dtypes(exclude=[np.number]).columns

numeric_transformer_pipeline = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])

categorical_transformer_pipeline = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer_pipeline, numeric_features),
        ('cat', categorical_transformer_pipeline, categorical_features)
    ],
    remainder='passthrough'
)

X_processed = preprocessor.fit_transform(X)

print("Preprocessing complete. Processed features shape:", X_processed.shape)

print("\nTarget variable (first 5 values):")
print(y.head())

```

Preprocessing complete. Processed features shape: (900, 7)

Target variable (first 5 values):

0	Kecimen
1	Kecimen
2	Kecimen
3	Kecimen
4	Kecimen

Name: Class, dtype: object

```

X_processed_df = pd.DataFrame(X_processed, columns=[f'processed_feature_{i}' for i in range(X_processed.shape[1])])

rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_processed_df, y)
importances_rf = rf.feature_importances_
feat_importance_rf = pd.Series(importances_rf, index=X_processed_df.columns).sort_values(ascending=False)
print("Random Forest Feature Importances (on processed data):")
print(feat_importance_rf.head(10))

k_best = SelectKBest(score_func=f_classif, k='all')
k_best.fit(X_processed_df, y)
kbest_feat_scores = pd.Series(k_best.scores_, index=X_processed_df.columns).sort_values(ascending=False)
print("SelectKBest Feature Scores (on processed data):")
print(kbest_feat_scores.head(10))

k_best_processed = 6
selected_features_kbest = kbest_feat_scores.head(k_best_processed).index.tolist()
print(f"\nTop {k_best_processed} features selected by SelectKBest: {selected_features_kbest}")

```

Random Forest Feature Importances (on processed data):

processed_feature_6	0.252325
processed_feature_1	0.244634
processed_feature_4	0.158066
processed_feature_0	0.124583
processed_feature_3	0.085957
processed_feature_2	0.070127
processed_feature_5	0.064307

dtype: float64

SelectKBest Feature Scores (on processed data):

processed_feature_1	744.251543
processed_feature_6	715.744169
processed_feature_0	577.808076
processed_feature_4	577.356620
processed_feature_2	304.321568
processed_feature_3	213.774500
processed_feature_5	22.014600

dtype: float64

Top 6 features selected by SelectKBest: ['processed_feature_1', 'processed_feature_6', 'processed_feature_0', 'processed_feature_4', 'processed_feature_2', 'processed_feature_3']

▼ Train Test Split

```

label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)

X_split = X_processed
y_split = y_encoded

X_train_80, X_test_20, y_train_80, y_test_20 = train_test_split(X_split, y_split, test_size=0.2, random_state=42, stratify=y)
print(f"80/20 split: Train shape: {X_train_80.shape}, Test shape: {X_test_20.shape}")

X_train_70, X_test_30, y_train_70, y_test_30 = train_test_split(X_split, y_split, test_size=0.3, random_state=42, stratify=y)
print(f"70/30 split: Train shape: {X_train_70.shape}, Test shape: {X_test_30.shape}")

X_train_60, X_test_40, y_train_60, y_test_40 = train_test_split(X_split, y_split, test_size=0.4, random_state=42, stratify=y)
print(f"60/40 split: Train shape: {X_train_60.shape}, Test shape: {X_test_40.shape}")

80/20 split: Train shape: (720, 7), Test shape: (180, 7)
70/30 split: Train shape: (630, 7), Test shape: (270, 7)
60/40 split: Train shape: (540, 7), Test shape: (360, 7)

```

We split the data into training and testing sets (80/20, 70/30, and 60/40) to evaluate model generalization.

- **Larger Training Set (e.g., 80/20):** Better for training complex models, but evaluation on the smaller test set might be less precise.
- **Smaller Training Set (e.g., 60/40):** Provides a more reliable performance estimate on the larger test set, but the model might underfit due to less training data.

Evaluating across these splits helps understand the impact of training data volume on model performance and the confidence in the evaluation metric. Stratified splitting ensures class distribution is maintained in each split.

Model Building

```

print("Training SVM model with PCA on processed data (80/20 split)...")

svm_pipeline_pca = Pipeline(steps=[('pca', PCA(n_components=0.95, random_state=42)),
                                    ('classifier', SVC(random_state=42, probability=True))])

svm_pipeline_pca.fit(X_train_80, y_train_80)

y_pred_svm_pca = svm_pipeline_pca.predict(X_test_20)

print("SVM Accuracy (80/20 split) with PCA:", accuracy_score(y_test_20, y_pred_svm_pca))

n_classes = len(np.unique(y_train_80))

if n_classes > 2:
    print("\nClassification Report (with PCA):\n", classification_report(y_test_20, y_pred_svm_pca))
    print("\nConfusion Matrix (with PCA):\n", confusion_matrix(y_test_20, y_pred_svm_pca))
else:
    print("\nClassification Report (with PCA):\n", classification_report(y_test_20, y_pred_svm_pca, zero_division=0))
    print("\nConfusion Matrix (with PCA):\n", confusion_matrix(y_test_20, y_pred_svm_pca))

if n_classes == 2:
    y_prob_svm_pca = svm_pipeline_pca.predict_proba(X_test_20)[:, 1]
    fpr, tpr, thresholds = roc_curve(y_test_20, y_prob_svm_pca)
    roc_auc = auc(fpr, tpr)

    plt.figure(figsize=(8, 6))
    plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver Operating Characteristic (ROC) Curve - SVM with PCA')
    plt.legend(loc="lower right")
    plt.show()

```

Training SVM model with PCA on processed data (80/20 split)...
SVM Accuracy (80/20 split) with PCA: 0.8833333333333333


Classification Report (with PCA):

	precision	recall	f1-score	support
0	0.95	0.81	0.87	90
1	0.83	0.96	0.89	90
accuracy			0.88	180
macro avg	0.89	0.88	0.88	180
weighted avg	0.89	0.88	0.88	180

Confusion Matrix (with PCA):

```
[[73 17]
 [ 4 86]]
```

Receiver Operating Characteristic (ROC) Curve - SVM with PCA



SVM Hyperparameter Justification:

- `C` (Regularization parameter): Controls the trade-off between classification error and margin size.
- `gamma` (Kernel coefficient): Influences the shape of the decision boundary, especially with non-linear kernels.
- `kernel` (Kernel type): Determines the type of hyperplane used for separation ('rbf', 'linear').

```
param_grid = {
    'classifier__C': [0.1, 1, 10, 100],
    'classifier__gamma': ['scale', 'auto'],
    'classifier__kernel': ['rbf', 'linear']
}

svm_pipeline_tuned = Pipeline(steps=[('pca', PCA(n_components=0.95, random_state=42)),
                                     ('classifier', SVC(random_state=42, probability=True))])

print("Performing GridSearchCV with PCA (80/20 split)...")
grid_search_80 = GridSearchCV(svm_pipeline_tuned, param_grid, cv=3, scoring='accuracy', n_jobs=-1)
grid_search_80.fit(X_train_80, y_train_80)
print("\nBest parameters found by GridSearchCV (80/20 split with PCA):")
print(grid_search_80.best_params_)
print("Best accuracy found by GridSearchCV (80/20 split with PCA):", grid_search_80.best_score_)
y_pred_grid_80 = grid_search_80.predict(X_test_20)
print("\nGridSearchCV Test Accuracy (80/20 split) with PCA:", accuracy_score(y_test_20, y_pred_grid_80))

print("\nPerforming GridSearchCV with PCA (70/30 split)...")
grid_search_70 = GridSearchCV(svm_pipeline_tuned, param_grid, cv=3, scoring='accuracy', n_jobs=-1)
grid_search_70.fit(X_train_70, y_train_70)
print("\nBest parameters found by GridSearchCV (70/30 split with PCA):")
print(grid_search_70.best_params_)
print("Best accuracy found by GridSearchCV (70/30 split with PCA):", grid_search_70.best_score_)
y_pred_grid_70 = grid_search_70.predict(X_test_30)
print("\nGridSearchCV Test Accuracy (70/30 split) with PCA:", accuracy_score(y_test_30, y_pred_grid_70))

print("\nPerforming GridSearchCV with PCA (60/40 split)...")


https://colab.research.google.com/drive/1LdokWCnyuMXjnePwIPm7rLOS8E3OAPu2#scrollTo=3e65fa97&printMode=true
```

```

grid_search_60 = GridSearchCV(svm_pipeline_tuned, param_grid, cv=3, scoring='accuracy', n_jobs=-1)
grid_search_60.fit(X_train_60, y_train_60)
print("\nBest parameters found by GridSearchCV (60/40 split with PCA):")
print(grid_search_60.best_params_)
print("Best accuracy found by GridSearchCV (60/40 split with PCA):", grid_search_60.best_score_)
y_pred_grid_60 = grid_search_60.predict(X_test_40)
print("\nGridSearchCV Test Accuracy (60/40 split) with PCA:", accuracy_score(y_test_40, y_pred_grid_60))

n_classes = len(np.unique(y_train_80))
class_names = label_encoder.classes_

print("\n--- GridSearchCV Classification Reports and Confusion Matrices ---")

print("\nGridSearchCV Classification Report (80/20 split with PCA):\n", classification_report(y_test_20, y_pred_grid_80, t):
print("\nGridSearchCV Confusion Matrix (80/20 split with PCA):")
cm_grid_80 = confusion_matrix(y_test_20, y_pred_grid_80)
ConfusionMatrixDisplay(cm_grid_80, display_labels=class_names).plot(cmap='Blues')
plt.title('GridSearchCV SVM with PCA - Confusion Matrix (80/20 Split)')
plt.show()

print("\nGridSearchCV Classification Report (70/30 split with PCA):\n", classification_report(y_test_30, y_pred_grid_70, t:
print("\nGridSearchCV Confusion Matrix (70/30 split with PCA):")
cm_grid_70 = confusion_matrix(y_test_30, y_pred_grid_70)
ConfusionMatrixDisplay(cm_grid_70, display_labels=class_names).plot(cmap='Blues')
plt.title('GridSearchCV SVM with PCA - Confusion Matrix (70/30 Split)')
plt.show()

print("\nGridSearchCV Classification Report (60/40 split with PCA):\n", classification_report(y_test_40, y_pred_grid_60, t:
print("\nGridSearchCV Confusion Matrix (60/40 split with PCA):")
cm_grid_60 = confusion_matrix(y_test_40, y_pred_grid_60)
ConfusionMatrixDisplay(cm_grid_60, display_labels=class_names).plot(cmap='Blues')
plt.title('GridSearchCV SVM with PCA - Confusion Matrix (60/40 Split)')
plt.show()

param_dist = {
    'classifier__C': np.logspace(-2, 2, 50),
    'classifier__gamma': ['scale', 'auto'],
    'classifier__kernel': ['rbf', 'linear']
}

print("\nPerforming RandomizedSearchCV with PCA (80/20 split)...")
random_search_80 = RandomizedSearchCV(svm_pipeline_tuned, param_dist, n_iter=10, cv=3, scoring='accuracy', random_state=42,
random_search_80.fit(X_train_80, y_train_80)
print("\nBest parameters found by RandomizedSearchCV (80/20 split with PCA):")
print(random_search_80.best_params_)
print("Best accuracy found by RandomizedSearchCV (80/20 split with PCA):", random_search_80.best_score_)
y_pred_random_80 = random_search_80.predict(X_test_20)
print("\nRandomizedSearchCV Test Accuracy (80/20 split) with PCA:", accuracy_score(y_test_20, y_pred_random_80))

print("\nPerforming RandomizedSearchCV with PCA (70/30 split)...")
random_search_70 = RandomizedSearchCV(svm_pipeline_tuned, param_dist, n_iter=10, cv=3, scoring='accuracy', random_state=42,
random_search_70.fit(X_train_70, y_train_70)
print("\nBest parameters found by RandomizedSearchCV (70/30 split with PCA):")
print(random_search_70.best_params_)
print("Best accuracy found by RandomizedSearchCV (70/30 split with PCA):", random_search_70.best_score_)
y_pred_random_70 = random_search_70.predict(X_test_30)
print("\nRandomizedSearchCV Test Accuracy (70/30 split) with PCA:", accuracy_score(y_test_30, y_pred_random_70))

print("\nPerforming RandomizedSearchCV with PCA (60/40 split)...")
random_search_60 = RandomizedSearchCV(svm_pipeline_tuned, param_dist, n_iter=10, cv=3, scoring='accuracy', random_state=42,
random_search_60.fit(X_train_60, y_train_60)
print("\nBest parameters found by RandomizedSearchCV (60/40 split with PCA):")
print(random_search_60.best_params_)
print("Best accuracy found by RandomizedSearchCV (60/40 split with PCA):", random_search_60.best_score_)
y_pred_random_60 = random_search_60.predict(X_test_40)
print("\nRandomizedSearchCV Test Accuracy (60/40 split) with PCA:", accuracy_score(y_test_40, y_pred_random_60))

print("\n--- RandomizedSearchCV Classification Reports and Confusion Matrices ---")

print("\nRandomizedSearchCV Classification Report (80/20 split with PCA):\n", classification_report(y_test_20, y_pred_random_80, t):
print("\nRandomizedSearchCV Confusion Matrix (80/20 split with PCA):")
cm_random_80 = confusion_matrix(y_test_20, y_pred_random_80)
ConfusionMatrixDisplay(cm_random_80, display_labels=class_names).plot(cmap='Blues')
plt.title('RandomizedSearchCV SVM with PCA - Confusion Matrix (80/20 Split)')
plt.show()

print("\nRandomizedSearchCV Classification Report (70/30 split with PCA):\n", classification_report(y_test_30, y_pred_random_70, t:
print("\nRandomizedSearchCV Confusion Matrix (70/30 split with PCA):")
cm_random_70 = confusion_matrix(y_test_30, y_pred_random_70)
ConfusionMatrixDisplay(cm_random_70, display_labels=class_names).plot(cmap='Blues')

```

```

plt.title('RandomizedSearchCV SVM with PCA - Confusion Matrix (70/30 Split)')
plt.show()

print("\nRandomizedSearchCV Classification Report (60/40 split with PCA):\n", classification_report(y_test_40, y_pred_random))
print("\nRandomizedSearchCV Confusion Matrix (60/40 split with PCA):")
cm_random_60 = confusion_matrix(y_test_40, y_pred_random_60)
ConfusionMatrixDisplay(cm_random_60, display_labels=class_names).plot(cmap='Blues')
plt.title('RandomizedSearchCV SVM with PCA - Confusion Matrix (60/40 Split)')
plt.show()

if n_classes == 2:
    print("\n--- GridSearchCV ROC Curves ---")
    plt.figure(figsize=(8, 6))

    # 80/20 split ROC
    y_prob_grid_80 = grid_search_80.best_estimator_.predict_proba(X_test_20)[:, 1]
    fpr_grid_80, tpr_grid_80, _ = roc_curve(y_test_20, y_prob_grid_80)
    roc_auc_grid_80 = auc(fpr_grid_80, tpr_grid_80)
    plt.plot(fpr_grid_80, tpr_grid_80, color='darkorange', lw=2, label='GridSearchCV 80/20 (area = %.2f)' % roc_auc_grid_80)

    # 70/30 split ROC
    y_prob_grid_70 = grid_search_70.best_estimator_.predict_proba(X_test_30)[:, 1]
    fpr_grid_70, tpr_grid_70, _ = roc_curve(y_test_30, y_prob_grid_70)
    roc_auc_grid_70 = auc(fpr_grid_70, tpr_grid_70)
    plt.plot(fpr_grid_70, tpr_grid_70, color='green', lw=2, label='GridSearchCV 70/30 (area = %.2f)' % roc_auc_grid_70)

    # 60/40 split ROC
    y_prob_grid_60 = grid_search_60.best_estimator_.predict_proba(X_test_40)[:, 1]
    fpr_grid_60, tpr_grid_60, _ = roc_curve(y_test_40, y_prob_grid_60)
    roc_auc_grid_60 = auc(fpr_grid_60, tpr_grid_60)
    plt.plot(fpr_grid_60, tpr_grid_60, color='red', lw=2, label='GridSearchCV 60/40 (area = %.2f)' % roc_auc_grid_60)

    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver Operating Characteristic (ROC) Curve - GridSearchCV SVM with PCA')
    plt.legend(loc="lower right")
    plt.show()

print("\n--- RandomizedSearchCV ROC Curves ---")
plt.figure(figsize=(8, 6))

# 80/20 split ROC
y_prob_random_80 = random_search_80.best_estimator_.predict_proba(X_test_20)[:, 1]
fpr_random_80, tpr_random_80, _ = roc_curve(y_test_20, y_prob_random_80)
roc_auc_random_80 = auc(fpr_random_80, tpr_random_80)
plt.plot(fpr_random_80, tpr_random_80, color='darkorange', lw=2, label='RandomizedSearchCV 80/20 (area = %.2f)' % roc_auc_random_80)

# 70/30 split ROC
y_prob_random_70 = random_search_70.best_estimator_.predict_proba(X_test_30)[:, 1]
fpr_random_70, tpr_random_70, _ = roc_curve(y_test_30, y_prob_random_70)
roc_auc_random_70 = auc(fpr_random_70, tpr_random_70)
plt.plot(fpr_random_70, tpr_random_70, color='green', lw=2, label='RandomizedSearchCV 70/30 (area = %.2f)' % roc_auc_random_70)

# 60/40 split ROC
y_prob_random_60 = random_search_60.best_estimator_.predict_proba(X_test_40)[:, 1]
fpr_random_60, tpr_random_60, _ = roc_curve(y_test_40, y_prob_random_60)
roc_auc_random_60 = auc(fpr_random_60, tpr_random_60)
plt.plot(fpr_random_60, tpr_random_60, color='red', lw=2, label='RandomizedSearchCV 60/40 (area = %.2f)' % roc_auc_random_60)

plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve - RandomizedSearchCV SVM with PCA')
plt.legend(loc="lower right")
plt.show()

```


Performing GridSearchCV with PCA (80/20 split)...

Best parameters found by GridSearchCV (80/20 split with PCA):
 {'classifier__C': 0.1, 'classifier__gamma': 'scale', 'classifier__kernel': 'rbf'}
 Best accuracy found by GridSearchCV (80/20 split with PCA): 0.8666666666666667

GridSearchCV Test Accuracy (80/20 split) with PCA: 0.8722222222222222

Performing GridSearchCV with PCA (70/30 split)...

Best parameters found by GridSearchCV (70/30 split with PCA):
 {'classifier__C': 0.1, 'classifier__gamma': 'scale', 'classifier__kernel': 'rbf'}
 Best accuracy found by GridSearchCV (70/30 split with PCA): 0.8793650793650793

GridSearchCV Test Accuracy (70/30 split) with PCA: 0.837037037037037

Performing GridSearchCV with PCA (60/40 split)...

Best parameters found by GridSearchCV (60/40 split with PCA):
 {'classifier__C': 0.1, 'classifier__gamma': 'scale', 'classifier__kernel': 'rbf'}
 Best accuracy found by GridSearchCV (60/40 split with PCA): 0.8722222222222222


GridSearchCV Test Accuracy (60/40 split) with PCA: 0.8527777777777777

--- GridSearchCV Classification Reports and Confusion Matrices ---

GridSearchCV Classification Report (80/20 split with PCA):

	precision	recall	f1-score	support
Besni	0.94	0.80	0.86	90
Kecimen	0.83	0.94	0.88	90
accuracy			0.87	180
macro avg	0.88	0.87	0.87	180
weighted avg	0.88	0.87	0.87	180


GridSearchCV Confusion Matrix (80/20 split with PCA):




GridSearchCV Classification Report (70/30 split with PCA):

	precision	recall	f1-score	support
Besni	0.89	0.77	0.83	135
Kecimen	0.80	0.90	0.85	135
accuracy			0.84	270
macro avg	0.84	0.84	0.84	270
weighted avg	0.84	0.84	0.84	270

GridSearchCV Confusion Matrix (70/30 split with PCA):





GridSearchCV Classification Report (60/40 split with PCA):


precision recall f1-score support

	Besni	0.90	0.79	0.84	180
	Kecimen	0.82	0.91	0.86	180

	accuracy	0.85	0.85	0.85	360
	macro avg	0.86	0.85	0.85	360
	weighted avg	0.86	0.85	0.85	360

GridSearchCV Confusion Matrix (60/40 split with PCA):

GridSearchCV SVM with PCA - Confusion Matrix (60/40 Split)



Performing RandomizedSearchCV with PCA (80/20 split)...

Best parameters found by RandomizedSearchCV (80/20 split with PCA):

```
{'classifier__kernel': 'linear', 'classifier__gamma': 'scale', 'classifier__C': np.float64(0.2442053094548651)}
```

Best accuracy found by RandomizedSearchCV (80/20 split with PCA): 0.8666666666666666

RandomizedSearchCV Test Accuracy (80/20 split) with PCA: 0.8777777777777778

Performing RandomizedSearchCV with PCA (70/30 split)...

Best parameters found by RandomizedSearchCV (70/30 split with PCA):

```
{'classifier__kernel': 'linear', 'classifier__gamma': 'auto', 'classifier__C': np.float64(0.017575106248547922)}
```

Best accuracy found by RandomizedSearchCV (70/30 split with PCA): 0.8809523809523809

RandomizedSearchCV Test Accuracy (70/30 split) with PCA: 0.8444444444444444

Performing RandomizedSearchCV with PCA (60/40 split)...

Best parameters found by RandomizedSearchCV (60/40 split with PCA):

```
{'classifier__kernel': 'linear', 'classifier__gamma': 'auto', 'classifier__C': np.float64(0.017575106248547922)}
```

Best accuracy found by RandomizedSearchCV (60/40 split with PCA): 0.8685185185185186

RandomizedSearchCV Test Accuracy (60/40 split) with PCA: 0.8583333333333333

--- RandomizedSearchCV Classification Reports and Confusion Matrices ---


RandomizedSearchCV Classification Report (80/20 split with PCA):

precision recall f1-score support

	Besni	0.93	0.82	0.87	90
	Kecimen	0.84	0.93	0.88	90

	accuracy	0.88	0.88	0.88	180
	macro avg	0.88	0.88	0.88	180
	weighted avg	0.88	0.88	0.88	180


RandomizedSearchCV Confusion Matrix (80/20 split with PCA):



RandomizedSearchCV Classification Report (70/30 split with PCA):

	precision	recall	f1-score	support
Besni	0.88	0.79	0.84	135
Kecimen	0.81	0.90	0.85	135
accuracy			0.84	270
macro avg	0.85	0.84	0.84	270
weighted avg	0.85	0.84	0.84	270


RandomizedSearchCV Confusion Matrix (70/30 split with PCA):



RandomizedSearchCV Classification Report (60/40 split with PCA):

	precision	recall	f1-score	support
Besni	0.90	0.81	0.85	180
Kecimen	0.83	0.91	0.86	180
accuracy			0.86	360
macro avg	0.86	0.86	0.86	360
weighted avg	0.86	0.86	0.86	360

RandomizedSearchCV Confusion Matrix (60/40 split with PCA):



```

print("Training Logistic Regression model...")

lr_pipeline = Pipeline(steps=[('classifier', LogisticRegression(random_state=42, max_iter=1000))])

lr_pipeline.fit(X_train_80, y_train_80)

y_pred_lr = lr_pipeline.predict(X_test_20)

print("\nLogistic Regression Accuracy (80/20 split):", accuracy_score(y_test_20, y_pred_lr))


n_classes = len(np.unique(y_train_80))

if n_classes > 2:
    print("\nClassification Report (Logistic Regression):\n", classification_report(y_test_20, y_pred_lr))
    print("\nConfusion Matrix (Logistic Regression):\n", confusion_matrix(y_test_20, y_pred_lr))
else:
    print("\nClassification Report (Logistic Regression):\n", classification_report(y_test_20, y_pred_lr, zero_division=0))
    print("\nConfusion Matrix (Logistic Regression):\n", confusion_matrix(y_test_20, y_pred_lr))

if n_classes == 2:
    y_prob_lr = lr_pipeline.predict_proba(X_test_20)[:, 1]
    fpr, tpr, thresholds = roc_curve(y_test_20, y_prob_lr)
    roc_auc = auc(fpr, tpr)

    plt.figure(figsize=(8, 6))
    plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver Operating Characteristic (ROC) Curve - Logistic Regression')
    plt.legend(loc="lower right")
    plt.show()

```




--- RandomizedSearchCV ROC Curves ---

Confusion Matrix (Logistic Regression):

```

[[75 15]
 [5 85]]

```



Logistic Regression Hyperparameter Justification:

- **C** (Inverse of regularization strength): Controls the amount of regularization applied.
- **penalty** (Regularization type): Specifies the norm used for penalization ('l1', 'l2', 'elasticnet', 'none').
- **solver** (Optimization algorithm): Selects the algorithm for finding the optimal model parameters.
- **l1_ratio** (Elastic-Net mixing parameter): Balances L1 and L2 penalties when **penalty** is 'elasticnet'.

```

class_names = label_encoder.classes_

lr = LogisticRegression(random_state=42, max_iter=2000)

print("\n--- Starting GridSearchCV for LogisticRegression (80/20 split) ---")

param_grid = [
    {
        'solver': ['lbfgs'],
        'penalty': ['l2', None],
        'C': [0.01, 0.1, 1, 10, 100]
    },
    {
        'solver': ['liblinear'],
        'penalty': ['l1', 'l2'],
        'C': [0.01, 0.1, 1, 10, 100]
    },
    {
        'solver': ['saga'],
        'penalty': ['l1', 'l2', 'elasticnet'],
        'C': [0.01, 0.1, 1, 10, 100],
        'l1_ratio': [0.5]
    }
]

grid_search_lr_80 = GridSearchCV(
    estimator=lr,
    param_grid=param_grid,
    cv=3,
    scoring='accuracy',
    n_jobs=-1,
    verbose=1
)

grid_search_lr_80.fit(X_train_80, y_train_80)

print("\n--- GridSearchCV (LogisticRegression) Complete (80/20 split) ---")
print(f"Best parameters found (80/20 split): {grid_search_lr_80.best_params_}")
print(f"Best cross-validation accuracy (80/20 split): {grid_search_lr_80.best_score_.:.4f}")

y_pred_grid_lr_80 = grid_search_lr_80.predict(X_test_20)
print(f"\nGridSearchCV Test Accuracy (80/20 split): {accuracy_score(y_test_20, y_pred_grid_lr_80):.4f}")

# Evaluate on 70/30 split
print("\n--- Starting GridSearchCV for LogisticRegression (70/30 split) ---")
grid_search_lr_70 = GridSearchCV(
    estimator=lr,
    param_grid=param_grid,
    cv=3,
    scoring='accuracy',
    n_jobs=-1,
    verbose=1
)
# Assuming X_train_70, y_train_70, X_test_30, y_test_30 are already defined
grid_search_lr_70.fit(X_train_70, y_train_70)
print("\n--- GridSearchCV (LogisticRegression) Complete (70/30 split) ---")
print(f"Best parameters found (70/30 split): {grid_search_lr_70.best_params_}")
print(f"Best cross-validation accuracy (70/30 split): {grid_search_lr_70.best_score_.:.4f}")
y_pred_grid_lr_70 = grid_search_lr_70.predict(X_test_30)
print(f"\nGridSearchCV Test Accuracy (70/30 split): {accuracy_score(y_test_30, y_pred_grid_lr_70):.4f}")

# Evaluate on 60/40 split
print("\n--- Starting GridSearchCV for LogisticRegression (60/40 split) ---")
grid_search_lr_60 = GridSearchCV(
    estimator=lr,
    param_grid=param_grid,
    cv=3,
    scoring='accuracy',
    n_jobs=-1,
    verbose=1
)

```

```

# Assuming X_train_60, y_train_60, X_test_40, y_test_40 are already defined
grid_search_lr_60.fit(X_train_60, y_train_60)
print("\n--- GridSearchCV (LogisticRegression) Complete (60/40 split) ---")
print(f"Best parameters found (60/40 split): {grid_search_lr_60.best_params_}")
print(f"Best cross-validation accuracy (60/40 split): {grid_search_lr_60.best_score_:.4f}")
y_pred_grid_lr_60 = grid_search_lr_60.predict(X_test_40)
print(f"\nGridSearchCV Test Accuracy (60/40 split): {accuracy_score(y_test_40, y_pred_grid_lr_60):.4f}")


if n_classes > 2:
    print("\nGridSearchCV Classification Report (80/20 split):\n", classification_report(y_test_20, y_pred_grid_lr_80, target_names=class_names))
    print("\nGridSearchCV Confusion Matrix (80/20 split):\n", confusion_matrix(y_test_20, y_pred_grid_lr_80))

    print("\nGridSearchCV Classification Report (70/30 split):\n", classification_report(y_test_30, y_pred_grid_lr_70, target_names=class_names))
    print("\nGridSearchCV Confusion Matrix (70/30 split):\n", confusion_matrix(y_test_30, y_pred_grid_lr_70))

    print("\nGridSearchCV Classification Report (60/40 split):\n", classification_report(y_test_40, y_pred_grid_lr_60, target_names=class_names))
    print("\nGridSearchCV Confusion Matrix (60/40 split):\n", confusion_matrix(y_test_40, y_pred_grid_lr_60))
else:
    print("\nGridSearchCV Classification Report (80/20 split):\n", classification_report(y_test_20, y_pred_grid_lr_80, target_names=class_names))
    print("\nGridSearchCV Confusion Matrix (80/20 split):\n", confusion_matrix(y_test_20, y_pred_grid_lr_80))
    cm_grid_80 = confusion_matrix(y_test_20, y_pred_grid_lr_80)
    ConfusionMatrixDisplay(cm_grid_80, display_labels=class_names).plot(cmap='Blues')
    plt.title('GridSearchCV Logistic Regression - Confusion Matrix (80/20 Split)')
    plt.show()

    print("\nGridSearchCV Classification Report (70/30 split):\n", classification_report(y_test_30, y_pred_grid_lr_70, target_names=class_names))
    print("\nGridSearchCV Confusion Matrix (70/30 split):\n", confusion_matrix(y_test_30, y_pred_grid_lr_70))
    cm_grid_70 = confusion_matrix(y_test_30, y_pred_grid_lr_70)
    ConfusionMatrixDisplay(cm_grid_70, display_labels=class_names).plot(cmap='Blues')
    plt.title('GridSearchCV Logistic Regression - Confusion Matrix (70/30 Split)')
    plt.show()

    print("\nGridSearchCV Classification Report (60/40 split):\n", classification_report(y_test_40, y_pred_grid_lr_60, target_names=class_names))
    print("\nGridSearchCV Confusion Matrix (60/40 split):\n", confusion_matrix(y_test_40, y_pred_grid_lr_60))
    cm_grid_60 = confusion_matrix(y_test_40, y_pred_grid_lr_60)
    ConfusionMatrixDisplay(cm_grid_60, display_labels=class_names).plot(cmap='Blues')
    plt.title('GridSearchCV Logistic Regression - Confusion Matrix (60/40 Split)')
    plt.show()


print("\n--- Starting RandomizedSearchCV for LogisticRegression (80/20 split) ---")

param_dist = [
    {
        'solver': ['lbfgs'],
        'penalty': ['l2', None],
        'C': loguniform(0.01, 100)
    },
    {
        'solver': ['liblinear'],
        'penalty': ['l1', 'l2'],
        'C': loguniform(0.01, 100)
    },
    {
        'solver': ['saga'],
        'penalty': ['l1', 'l2', 'elasticnet'],
        'C': loguniform(0.01, 100),
        'l1_ratio': np.linspace(0, 1, 10) # Samples from a list
    }
]

random_search_lr_80 = RandomizedSearchCV(
    estimator=lr,
    param_distributions=param_dist,
    n_iter=10,
    cv=3,
    scoring='accuracy',
    n_jobs=-1,
    random_state=42,
    verbose=1
)

random_search_lr_80.fit(X_train_80, y_train_80)

print("\n--- RandomizedSearchCV (LogisticRegression) Complete (80/20 split) ---")
print(f"Best parameters found (80/20 split): {random_search_lr_80.best_params_}")
print(f"Best cross-validation accuracy (80/20 split): {random_search_lr_80.best_score_:.4f}")

```

```

y_pred_random_lr_80 = random_search_lr_80.predict(X_test_20)
print(f"\nRandomizedSearchCV Test Accuracy (80/20 split): {accuracy_score(y_test_20, y_pred_random_lr_80):.4f}")

# Evaluate on 70/30 split
print("\n--- Starting RandomizedSearchCV for LogisticRegression (70/30 split) ---")
random_search_lr_70 = RandomizedSearchCV(
    estimator=lr,
    param_distributions=param_dist,
    n_iter=10,
    cv=3,
    scoring='accuracy',
    n_jobs=-1,
    random_state=42,
    verbose=1
)
# Assuming X_train_70, y_train_70, X_test_30, y_test_30 are already defined
random_search_lr_70.fit(X_train_70, y_train_70)
print("\n--- RandomizedSearchCV (LogisticRegression) Complete (70/30 split) ---")
print(f"Best parameters found (70/30 split): {random_search_lr_70.best_params_}")
print(f"Best cross-validation accuracy (70/30 split): {random_search_lr_70.best_score_:.4f}")
y_pred_random_lr_70 = random_search_lr_70.predict(X_test_30)
print(f"\nRandomizedSearchCV Test Accuracy (70/30 split): {accuracy_score(y_test_30, y_pred_random_lr_70):.4f}")

# Evaluate on 60/40 split
print("\n--- Starting RandomizedSearchCV for LogisticRegression (60/40 split) ---")
random_search_lr_60 = RandomizedSearchCV(
    estimator=lr,
    param_distributions=param_dist,
    n_iter=10,
    cv=3,
    scoring='accuracy',
    n_jobs=-1,
    random_state=42,
    verbose=1
)
# Assuming X_train_60, y_train_60, X_test_40, y_test_40 are already defined
random_search_lr_60.fit(X_train_60, y_train_60)
print("\n--- RandomizedSearchCV (LogisticRegression) Complete (60/40 split) ---")
print(f"Best parameters found (60/40 split): {random_search_lr_60.best_params_}")
print(f"Best cross-validation accuracy (60/40 split): {random_search_lr_60.best_score_:.4f}")
y_pred_random_lr_60 = random_search_lr_60.predict(X_test_40)
print(f"\nRandomizedSearchCV Test Accuracy (60/40 split): {accuracy_score(y_test_40, y_pred_random_lr_60):.4f}")

# Classification Reports and Confusion Matrices for RandomizedSearchCV
print("\n--- RandomizedSearchCV Classification Reports and Confusion Matrices ---")

if n_classes > 2:
    print("\nRandomizedSearchCV Classification Report (80/20 split):\n", classification_report(y_test_20, y_pred_random_lr_80))
    print("\nRandomizedSearchCV Confusion Matrix (80/20 split):\n", confusion_matrix(y_test_20, y_pred_random_lr_80))

    print("\nRandomizedSearchCV Classification Report (70/30 split):\n", classification_report(y_test_30, y_pred_random_lr_70))
    print("\nRandomizedSearchCV Confusion Matrix (70/30 split):\n", confusion_matrix(y_test_30, y_pred_random_lr_70))

    print("\nRandomizedSearchCV Classification Report (60/40 split):\n", classification_report(y_test_40, y_pred_random_lr_60))
    print("\nRandomizedSearchCV Confusion Matrix (60/40 split):\n", confusion_matrix(y_test_40, y_pred_random_lr_60))
else:
    print("\nRandomizedSearchCV Classification Report (80/20 split):\n", classification_report(y_test_20, y_pred_random_lr_80))
    print("\nRandomizedSearchCV Confusion Matrix (80/20 split):\n", confusion_matrix(y_test_20, y_pred_random_lr_80))
    cm_random_80 = confusion_matrix(y_test_20, y_pred_random_lr_80)
    ConfusionMatrixDisplay(cm_random_80, display_labels=class_names).plot(cmap='Blues')
    plt.title('RandomizedSearchCV Logistic Regression - Confusion Matrix (80/20 Split)')
    plt.show()

    print("\nRandomizedSearchCV Classification Report (70/30 split):\n", classification_report(y_test_30, y_pred_random_lr_70))
    print("\nRandomizedSearchCV Confusion Matrix (70/30 split):\n", confusion_matrix(y_test_30, y_pred_random_lr_70))
    cm_random_70 = confusion_matrix(y_test_30, y_pred_random_lr_70)
    ConfusionMatrixDisplay(cm_random_70, display_labels=class_names).plot(cmap='Blues')
    plt.title('RandomizedSearchCV Logistic Regression - Confusion Matrix (70/30 Split)')
    plt.show()

    print("\nRandomizedSearchCV Classification Report (60/40 split):\n", classification_report(y_test_40, y_pred_random_lr_60))
    print("\nRandomizedSearchCV Confusion Matrix (60/40 split):\n", confusion_matrix(y_test_40, y_pred_random_lr_60))
    cm_random_60 = confusion_matrix(y_test_40, y_pred_random_lr_60)
    ConfusionMatrixDisplay(cm_random_60, display_labels=class_names).plot(cmap='Blues')
    plt.title('RandomizedSearchCV Logistic Regression - Confusion Matrix (60/40 Split)')
    plt.show()

# ROC Curves for binary classification (Combined plots)
if n_classes == 2:
    print("\n--- GridSearchCV ROC Curves ---")

```

```
plt.figure(figsize=(8, 6))

# 80/20 split ROC
y_prob_grid_80 = grid_search_lr_80.best_estimator_.predict_proba(X_test_20)[:, 1]
fpr_grid_80, tpr_grid_80, _ = roc_curve(y_test_20, y_prob_grid_80)
roc_auc_grid_80 = auc(fpr_grid_80, tpr_grid_80)
plt.plot(fpr_grid_80, tpr_grid_80, color='darkorange', lw=2, label='GridSearchCV 80/20 (area = %0.2f)' % roc_auc_grid_80

# 70/30 split ROC
y_prob_grid_70 = grid_search_lr_70.best_estimator_.predict_proba(X_test_30)[:, 1]
fpr_grid_70, tpr_grid_70, _ = roc_curve(y_test_30, y_prob_grid_70)
roc_auc_grid_70 = auc(fpr_grid_70, tpr_grid_70)
plt.plot(fpr_grid_70, tpr_grid_70, color='green', lw=2, label='GridSearchCV 70/30 (area = %0.2f)' % roc_auc_grid_70

# 60/40 split ROC
y_prob_grid_60 = grid_search_lr_60.best_estimator_.predict_proba(X_test_40)[:, 1]
fpr_grid_60, tpr_grid_60, _ = roc_curve(y_test_40, y_prob_grid_60)
roc_auc_grid_60 = auc(fpr_grid_60, tpr_grid_60)
plt.plot(fpr_grid_60, tpr_grid_60, color='red', lw=2, label='GridSearchCV 60/40 (area = %0.2f)' % roc_auc_grid_60

plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve - GridSearchCV Logistic Regression')
plt.legend(loc="lower right")
plt.show()

print("\n--- RandomizedSearchCV ROC Curves ---")
plt.figure(figsize=(8, 6))

# 80/20 split ROC
y_prob_random_80 = random_search_lr_80.best_estimator_.predict_proba(X_test_20)[:, 1]
fpr_random_80, tpr_random_80, _ = roc_curve(y_test_20, y_prob_random_80)
roc_auc_random_80 = auc(fpr_random_80, tpr_random_80)
plt.plot(fpr_random_80, tpr_random_80, color='darkorange', lw=2, label='RandomizedSearchCV 80/20 (area = %0.2f)' % roc_auc_random_80

# 70/30 split ROC
y_prob_random_70 = random_search_lr_70.best_estimator_.predict_proba(X_test_30)[:, 1]
fpr_random_70, tpr_random_70, _ = roc_curve(y_test_30, y_prob_random_70)
roc_auc_random_70 = auc(fpr_random_70, tpr_random_70)
plt.plot(fpr_random_70, tpr_random_70, color='green', lw=2, label='RandomizedSearchCV 70/30 (area = %0.2f)' % roc_auc_random_70

# 60/40 split ROC
y_prob_random_60 = random_search_lr_60.best_estimator_.predict_proba(X_test_40)[:, 1]
fpr_random_60, tpr_random_60, _ = roc_curve(y_test_40, y_prob_random_60)
roc_auc_random_60 = auc(fpr_random_60, tpr_random_60)
plt.plot(fpr_random_60, tpr_random_60, color='red', lw=2, label='RandomizedSearchCV 60/40 (area = %0.2f)' % roc_auc_random_60

plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve - RandomizedSearchCV Logistic Regression')
plt.legend(loc="lower right")
plt.show()
```



```

--- Starting GridSearchCV for LogisticRegression (80/20 split) ---
Fitting 3 folds for each of 35 candidates, totalling 105 fits
/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logistic.py:1196: UserWarning: l1_ratio parameter is only us
  warnings.warn(
/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_sag.py:348: ConvergenceWarning: The max_iter was reached whi
  warnings.warn(
...  

--- GridSearchCV (LogisticRegression) Complete (80/20 split) ---
Best parameters found (80/20 split): {'C': 100, 'l1_ratio': 0.5, 'penalty': 'l1', 'solver': 'saga'}
Best cross-validation accuracy (80/20 split): 0.8681

GridSearchCV Test Accuracy (80/20 split): 0.8944

--- Starting GridSearchCV for LogisticRegression (70/30 split) ---
Fitting 3 folds for each of 35 candidates, totalling 105 fits
...  

--- GridSearchCV (LogisticRegression) Complete (70/30 split) ---
Best parameters found (70/30 split): {'C': 0.1, 'penalty': 'l2', 'solver': 'liblinear'}
Best cross-validation accuracy (70/30 split): 0.8825

GridSearchCV Test Accuracy (70/30 split): 0.8444

--- Starting GridSearchCV for LogisticRegression (60/40 split) ---
Fitting 3 folds for each of 35 candidates, totalling 105 fits
...  

--- GridSearchCV (LogisticRegression) Complete (60/40 split) ---
Best parameters found (60/40 split): {'C': 0.01, 'penalty': 'l2', 'solver': 'liblinear'}
Best cross-validation accuracy (60/40 split): 0.8704

GridSearchCV Test Accuracy (60/40 split): 0.8500

GridSearchCV Classification Report (80/20 split):
precision    recall   f1-score   support
...  


  Besni       0.95     0.83     0.89      90
  Kecimen     0.85     0.96     0.90      90
  accuracy           0.89      180
  macro avg       0.90     0.89     0.89      180
  weighted avg     0.90     0.89     0.89      180

```

GridSearchCV Confusion Matrix (80/20 split):

```
[[75 15]
 [ 4 86]]
```

GridSearchCV Logistic Regression - Confusion Matrix (80/20 Split)



GridSearchCV Classification Report (70/30 split):
precision recall f1-score support


```
...  


  Besni       0.88     0.79     0.84      135
  Kecimen     0.81     0.90     0.85      135
  accuracy           0.84      270
  macro avg       0.85     0.84     0.84      270
  weighted avg     0.85     0.84     0.84      270
```

GridSearchCV Confusion Matrix (70/30 split):

```
[[107 28]
 [ 14 121]]
```

GridSearchCV Logistic Regression - Confusion Matrix (70/30 Split)






```
GridSearchCV Classification Report (60/40 split):
precision    recall   f1-score   support
      Besni      0.92      0.77      0.84      180
      Kecimen     0.80      0.93      0.86      180
      accuracy           0.85      360
      macro avg       0.86      0.85      0.85      360
  weighted avg       0.86      0.85      0.85      360
```

```
GridSearchCV Confusion Matrix (60/40 split):
[[138  42]
 [ 12 168]]
```

GridSearchCV Logistic Regression - Confusion Matrix (60/40 Split)



```
--- Starting RandomizedSearchCV for LogisticRegression (80/20 split) ---
Fitting 3 folds for each of 10 candidates, totalling 30 fits
/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logistic.py:1196: UserWarning: l1_ratio parameter is only us
  warnings.warn(
/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_sag.py:348: ConvergenceWarning: The max_iter was reached whi
  warnings.warn

--- RandomizedSearchCV (LogisticRegression) Complete (80/20 split) ---
Best parameters found (80/20 split): {'C': np.float64(15.352246941973492), 'l1_ratio': np.float64(0.7777777777777777), 'pe
Best cross-validation accuracy (80/20 split): 0.8667

RandomizedSearchCV Test Accuracy (80/20 split): 0.8944

--- Starting RandomizedSearchCV for LogisticRegression (70/30 split) ---
Fitting 3 folds for each of 10 candidates, totalling 30 fits

--- RandomizedSearchCV (LogisticRegression) Complete (70/30 split) ---
Best parameters found (70/30 split): {'C': np.float64(0.04207988669606638), 'penalty': 'l2', 'solver': 'lbfgs'}
Best cross-validation accuracy (70/30 split): 0.8810

RandomizedSearchCV Test Accuracy (70/30 split): 0.8444

--- Starting RandomizedSearchCV for LogisticRegression (60/40 split) ---
Fitting 3 folds for each of 10 candidates, totalling 30 fits
```

```
-- RandomizedSearchCV (LogisticRegression) Complete (60/40 split) --
Best parameters found (60/40 split): {'C': np.float64(0.010071984838809197), 'penalty': 'l2', 'solver': 'liblinear'}
Best cross-validation accuracy (60/40 split): 0.8704
```

RandomizedSearchCV Test Accuracy (60/40 split): 0.8500

-- RandomizedSearchCV Classification Reports and Confusion Matrices --

RandomizedSearchCV Classification Report (80/20 split):


	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Besni	0.95	0.83	0.89	90
Kecimen	0.85	0.96	0.90	90
accuracy			0.89	180
macro avg	0.90	0.89	0.89	180
weighted avg	0.90	0.89	0.89	180

RandomizedSearchCV Confusion Matrix (80/20 split):

```
[[75 15]
 [ 4 86]]
```

RandomizedSearchCV Logistic Regression - Confusion Matrix (80/20 Split)



RandomizedSearchCV Classification Report (70/30 split):


	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Besni	0.88	0.79	0.84	135
Kecimen	0.81	0.90	0.85	135
accuracy			0.84	270
macro avg	0.85	0.84	0.84	270
weighted avg	0.85	0.84	0.84	270

RandomizedSearchCV Confusion Matrix (70/30 split):

```
[[107 28]
 [ 14 121]]
```

RandomizedSearchCV Logistic Regression - Confusion Matrix (70/30 Split)



Predicted label

RandomizedSearchCV Classification Report (60/40 split):

precision recall f1-score support


Besni	0.92	0.77	0.84	180
Kecimen	0.80	0.93	0.86	180
accuracy			0.85	360
macro avg	0.86	0.85	0.85	360
weighted avg	0.86	0.85	0.85	360

RandomizedSearchCV Confusion Matrix (60/40 split):

[[138 42]


[12 168]]

RandomizedSearchCV Logistic Regression - Confusion Matrix (60/40 Split)




--- GridSearchCV ROC Curves ---

Receiver Operating Characteristic (ROC) Curve - GridSearchCV Logistic Regression



--- RandomizedSearchCV ROC Curves ---

Receiver Operating Characteristic (ROC) Curve - RandomizedSearchCV Logistic Regression



```
print("Training Decision Tree model...")  
  
dt_pipeline = Pipeline(steps=[('classifier', DecisionTreeClassifier(random_state=42))])  
  
dt_pipeline.fit(X_train_80, y_train_80)  
  
y_pred_dt = dt_pipeline.predict(X_test_20)  
  
print("\nDecision Tree Accuracy (80/20 split):", accuracy_score(y_test_20, y_pred_dt))  
  
n_classes = len(np.unique(y_train_80))  
  
if n_classes > 2:  
    print("\nClassification Report (Decision Tree):\n", classification_report(y_test_20, y_pred_dt))  
    print("\nConfusion Matrix (Decision Tree):\n", confusion_matrix(y_test_20, y_pred_dt))  
else:  
    print("\nClassification Report (Decision Tree):\n", classification_report(y_test_20, y_pred_dt, zero_division=0))  
    print("\nConfusion Matrix (Decision Tree):\n", confusion_matrix(y_test_20, y_pred_dt))  
  
if n_classes == 2:  
    try:  
        y_prob_dt = dt_pipeline.predict_proba(X_test_20)[:, 1]  
        fpr, tpr, thresholds = roc_curve(y_test_20, y_prob_dt)  
        roc_auc = auc(fpr, tpr)  
  
        plt.figure(figsize=(8, 6))  
        plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)  
        plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')  
        plt.xlim([0.0, 1.0])  
        plt.ylim([0.0, 1.05])  
        plt.xlabel('False Positive Rate')  
        plt.ylabel('True Positive Rate')  
        plt.title('Receiver Operating Characteristic (ROC) Curve - Decision Tree')  
        plt.legend(loc="lower right")  
        plt.show()  
    except AttributeError:  
        print("\nDecision Tree classifier does not support predict_proba for ROC curve plotting.")
```

Training Decision Tree model...

Decision Tree Accuracy (80/20 split): 0.8444444444444444


Classification Report (Decision Tree):

	precision	recall	f1-score	support
0	0.90	0.78	0.83	90
1	0.80	0.91	0.85	90
accuracy			0.84	180
macro avg	0.85	0.84	0.84	180
weighted avg	0.85	0.84	0.84	180

Confusion Matrix (Decision Tree):

```
[[70 20]
 [ 8 82]]
```

Receiver Operating Characteristic (ROC) Curve - Decision Tree



Decision Tree Hyperparameter Justification:

- `criterion` (Split quality measure): Determines how the quality of a split is measured ('gini', 'entropy').
- `max_depth` (Maximum tree depth): Limits the depth of the tree to prevent overfitting.
- `min_samples_split` (Minimum samples to split): Sets the minimum number of samples required to split a node.
- `min_samples_leaf` (Minimum samples per leaf): Specifies the minimum number of samples required in a leaf node.

```
dt = DecisionTreeClassifier(random_state=42)
param_grid = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [3, 5, 7, 9, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

print("Performing GridSearchCV (80/20 split)...")
grid_search_80 = GridSearchCV(estimator=dt, param_grid=param_grid, cv=5, scoring='accuracy', n_jobs=-1)
grid_search_80.fit(X_train_80, y_train_80)
best_dt_80 = grid_search_80.best_estimator_
print("Best Parameters (80/20 split):", grid_search_80.best_params_)
y_pred_dt_grid_80 = grid_search_80.predict(X_test_20)
print("\nGridSearchCV Test Accuracy (80/20 split):", accuracy_score(y_test_20, y_pred_dt_grid_80))

print("\nPerforming GridSearchCV (70/30 split)...")
grid_search_70 = GridSearchCV(estimator=dt, param_grid=param_grid, cv=5, scoring='accuracy', n_jobs=-1)
grid_search_70.fit(X_train_70, y_train_70)
best_dt_70 = grid_search_70.best_estimator_
print("Best Parameters (70/30 split):", grid_search_70.best_params_)
y_pred_dt_grid_70 = grid_search_70.predict(X_test_30)
print("\nGridSearchCV Test Accuracy (70/30 split):", accuracy_score(y_test_30, y_pred_dt_grid_70))

print("\nPerforming GridSearchCV (60/40 split)...")
grid_search_60 = GridSearchCV(estimator=dt, param_grid=param_grid, cv=5, scoring='accuracy', n_jobs=-1)
```

```

grid_search_60.fit(X_train_60, y_train_60)
best_dt_60 = grid_search_60.best_estimator_
print("Best Parameters (60/40 split):", grid_search_60.best_params_)
y_pred_dt_grid_60 = grid_search_60.predict(X_test_40)
print("\nGridSearchCV Test Accuracy (60/40 split):", accuracy_score(y_test_40, y_pred_dt_grid_60))

n_classes = len(np.unique(y_train_80))
class_names = label_encoder.classes_

print("\n--- GridSearchCV Classification Reports and Confusion Matrices ---")

print("\nGridSearchCV Classification Report (80/20 split Decision Tree):\n", classification_report(y_test_20, y_pred_dt_grid_80))
print("\nGridSearchCV Confusion Matrix (80/20 split Decision Tree):")
cm_grid_dt_80 = confusion_matrix(y_test_20, y_pred_dt_grid_80)
ConfusionMatrixDisplay(cm_grid_dt_80, display_labels=class_names).plot(cmap='Blues')
plt.title('GridSearchCV Decision Tree - Confusion Matrix (80/20 Split)')
plt.show()

print("\nGridSearchCV Classification Report (70/30 split Decision Tree):\n", classification_report(y_test_30, y_pred_dt_grid_70))
print("\nGridSearchCV Confusion Matrix (70/30 split Decision Tree):")
cm_grid_dt_70 = confusion_matrix(y_test_30, y_pred_dt_grid_70)
ConfusionMatrixDisplay(cm_grid_dt_70, display_labels=class_names).plot(cmap='Blues')
plt.title('GridSearchCV Decision Tree - Confusion Matrix (70/30 Split)')
plt.show()

print("\nGridSearchCV Classification Report (60/40 split Decision Tree):\n", classification_report(y_test_40, y_pred_dt_grid_60))
print("\nGridSearchCV Confusion Matrix (60/40 split Decision Tree):")
cm_grid_dt_60 = confusion_matrix(y_test_40, y_pred_dt_grid_60)
ConfusionMatrixDisplay(cm_grid_dt_60, display_labels=class_names).plot(cmap='Blues')
plt.title('GridSearchCV Decision Tree - Confusion Matrix (60/40 Split)')
plt.show()

if n_classes == 2:
    try:
        print("\n--- GridSearchCV ROC Curves ---")
        plt.figure(figsize=(8, 6))

        y_prob_dt_grid_80 = grid_search_80.predict_proba(X_test_20)[:, 1]
        fpr_grid_dt_80, tpr_grid_dt_80, _ = roc_curve(y_test_20, y_prob_dt_grid_80)
        roc_auc_grid_dt_80 = auc(fpr_grid_dt_80, tpr_grid_dt_80)
        plt.plot(fpr_grid_dt_80, tpr_grid_dt_80, color='darkorange', lw=2, label='GridSearchCV 80/20 (area = %0.2f)' % roc_auc_grid_dt_80)

        y_prob_dt_grid_70 = grid_search_70.predict_proba(X_test_30)[:, 1]
        fpr_grid_dt_70, tpr_grid_dt_70, _ = roc_curve(y_test_30, y_prob_dt_grid_70)
        roc_auc_grid_dt_70 = auc(fpr_grid_dt_70, tpr_grid_dt_70)
        plt.plot(fpr_grid_dt_70, tpr_grid_dt_70, color='green', lw=2, label='GridSearchCV 70/30 (area = %0.2f)' % roc_auc_grid_dt_70)

        y_prob_dt_grid_60 = grid_search_60.predict_proba(X_test_40)[:, 1]
        fpr_grid_dt_60, tpr_grid_dt_60, _ = roc_curve(y_test_40, y_prob_dt_grid_60)
        roc_auc_grid_dt_60 = auc(fpr_grid_dt_60, tpr_grid_dt_60)
        plt.plot(fpr_grid_dt_60, tpr_grid_dt_60, color='red', lw=2, label='GridSearchCV 60/40 (area = %0.2f)' % roc_auc_grid_dt_60)

        plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
        plt.xlim([0.0, 1.0])
        plt.ylim([0.0, 1.05])
        plt.xlabel('False Positive Rate')
        plt.ylabel('True Positive Rate')
        plt.title('Receiver Operating Characteristic (ROC) Curve - GridSearchCV Decision Tree')
        plt.legend(loc="lower right")
        plt.show()

    except AttributeError:
        print("\nDecision Tree classifier does not support predict_proba for ROC curve plotting with this configuration.")

```



```

Performing GridSearchCV (80/20 split)...
Best Parameters (80/20 split): {'criterion': 'gini', 'max_depth': 3, 'min_samples_leaf': 1, 'min_samples_split': 10}

GridSearchCV Test Accuracy (80/20 split): 0.8722222222222222

Performing GridSearchCV (70/30 split)...
Best Parameters (70/30 split): {'criterion': 'gini', 'max_depth': 3, 'min_samples_leaf': 1, 'min_samples_split': 10}

GridSearchCV Test Accuracy (70/30 split): 0.833333333333334

Performing GridSearchCV (60/40 split)...
Best Parameters (60/40 split): {'criterion': 'entropy', 'max_depth': 5, 'min_samples_leaf': 1, 'min_samples_split': 5}

GridSearchCV Test Accuracy (60/40 split): 0.8222222222222222


--- GridSearchCV Classification Reports and Confusion Matrices ---

GridSearchCV Classification Report (80/20 split Decision Tree):
precision    recall   f1-score  support
Besni        0.92     0.81     0.86      90
Kecimen      0.83     0.93     0.88      90

accuracy          0.87      180
macro avg       0.88     0.87     0.87      180
weighted avg    0.88     0.87     0.87      180

```

GridSearchCV Confusion Matrix (80/20 split Decision Tree):




```

GridSearchCV Classification Report (70/30 split Decision Tree):
precision    recall   f1-score  support
Besni        0.86     0.80     0.83      135
Kecimen      0.81     0.87     0.84      135

accuracy          0.83      270
macro avg       0.83     0.83     0.83      270
weighted avg    0.83     0.83     0.83      270

```

GridSearchCV Confusion Matrix (70/30 split Decision Tree):



```

print("Training KNN model...")

knn_pipeline = Pipeline(steps=[('classifier', KNeighborsClassifier())])

knn_pipeline.fit(X_train_80, y_train_80)

y_pred_knn = knn_pipeline.predict(X_test_20)

print("\nKNN Accuracy (80/20 split):", accuracy_score(y_test_20, y_pred_knn))

```

```



n_classes = len(np.unique(y_train_80))

if n_classes > 2:
    print("\nClassification Report (KNN):\n", classification_report(y_test_20, y_pred_knn))
    print("\nConfusion Matrix (KNN):\n", confusion_matrix(y_test_20, y_pred_knn))
else:
    print("\nClassification Report (KNN):\n", classification_report(y_test_20, y_pred_knn, zero_division=0))
    print("\nConfusion Matrix (KNN):\n", confusion_matrix(y_test_20, y_pred_knn))

if n_classes == 2:
    y_prob_knn = knn_pipeline.predict_proba(X_test_20)[:, 1]
    fpr, tpr, thresholds = roc_curve(y_test_20, y_prob_knn)
    roc_auc = auc(fpr, tpr)

    plt.figure(figsize=(8, 6))
    plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver Operating Characteristic (ROC) Curve - KNN')
    plt.legend(loc="lower right")
    plt.show()

```



KNN Hyperparameter Justification:

- `n_neighbors` (Number of neighbors): Defines the number of nearest data points considered for classification.
- `weights` (Weighting of neighbors): Determines how the contribution of neighbors is weighted ('uniform', 'distance').
- `metric` (Distance metric): Specifies how the distance between data points is calculated ('euclidean', 'manhattan', 'minkowski').

```

X_train_80, X_test_20, y_train_80, y_test_20 = train_test_split(X_split, y_split, test_size=0.2, random_state=42, stratify=y)
X_train_70, X_test_30, y_train_70, y_test_30 = train_test_split(X_split, y_split, test_size=0.3, random_state=42, stratify=y)
X_train_60, X_test_40, y_train_60, y_test_40 = train_test_split(X_split, y_split, test_size=0.4, random_state=42, stratify=y)

y_test_20_encoded = y_test_20
y_test_30_encoded = y_test_30
y_test_40_encoded = y_test_40

pipe = Pipeline([('knn', KNeighborsClassifier())])

param_grid = {
    'knn__n_neighbors': range(3, 16),
    'knn__weights': ['uniform', 'distance'],
    'knn__metric': ['euclidean', 'manhattan', 'minkowski']
}

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

print("Performing GridSearchCV on KNN (80/20 split)...")
grid_80 = GridSearchCV(pipe, param_grid, cv=cv, scoring='accuracy', n_jobs=-1)
grid_80.fit(X_train_80, y_train_80)
print('Best Parameters (80/20 split):', grid_80.best_params_)
best_knn_80 = grid_80.best_estimator_
y_pred_80 = best_knn_80.predict(X_test_20)
y_prob_80 = best_knn_80.predict_proba(X_test_20)[:, 1] if len(np.unique(y_split)) == 2 else None
print("Test Accuracy (80/20 split):", accuracy_score(y_test_20, y_pred_80))

print("\nPerforming GridSearchCV on KNN (70/30 split)...")
grid_70 = GridSearchCV(pipe, param_grid, cv=cv, scoring='accuracy', n_jobs=-1)
grid_70.fit(X_train_70, y_train_70)
print('Best Parameters (70/30 split):', grid_70.best_params_)
best_knn_70 = grid_70.best_estimator_
y_pred_70 = best_knn_70.predict(X_test_30)
y_prob_70 = best_knn_70.predict_proba(X_test_30)[:, 1] if len(np.unique(y_split)) == 2 else None
print("Test Accuracy (70/30 split):", accuracy_score(y_test_30, y_pred_70))

print("\nPerforming GridSearchCV on KNN (60/40 split)...")
grid_60 = GridSearchCV(pipe, param_grid, cv=cv, scoring='accuracy', n_jobs=-1)
grid_60.fit(X_train_60, y_train_60)
print('Best Parameters (60/40 split):', grid_60.best_params_)
best_knn_60 = grid_60.best_estimator_
y_pred_60 = best_knn_60.predict(X_test_40)
y_prob_60 = best_knn_60.predict_proba(X_test_40)[:, 1] if len(np.unique(y_split)) == 2 else None
print("Test Accuracy (60/40 split):", accuracy_score(y_test_40, y_pred_60))

class_names = label_encoder.classes_

print("\n--- GridSearchCV Classification Reports and Confusion Matrices ---")

print("\nGridSearchCV Classification Report (80/20 split KNN):\n", classification_report(y_test_20, y_pred_80, target_names=class_names))
print("\nGridSearchCV Confusion Matrix (80/20 split KNN):")
cm_grid_80 = confusion_matrix(y_test_20, y_pred_80)
ConfusionMatrixDisplay(cm_grid_80, display_labels=class_names).plot(cmap='Blues')
plt.title('GridSearchCV KNN - Confusion Matrix (80/20 Split)')
plt.show()

print("\nGridSearchCV Classification Report (70/30 split KNN):\n", classification_report(y_test_30, y_pred_70, target_names=class_names))
print("\nGridSearchCV Confusion Matrix (70/30 split KNN):")
cm_grid_70 = confusion_matrix(y_test_30, y_pred_70)
ConfusionMatrixDisplay(cm_grid_70, display_labels=class_names).plot(cmap='Blues')
plt.title('GridSearchCV KNN - Confusion Matrix (70/30 Split)')
plt.show()

print("\nGridSearchCV Classification Report (60/40 split KNN):\n", classification_report(y_test_40, y_pred_60, target_names=class_names))
print("\nGridSearchCV Confusion Matrix (60/40 split KNN):")
cm_grid_60 = confusion_matrix(y_test_40, y_pred_60)
ConfusionMatrixDisplay(cm_grid_60, display_labels=class_names).plot(cmap='Blues')
plt.title('GridSearchCV KNN - Confusion Matrix (60/40 Split)')
plt.show()

if len(np.unique(y_split)) == 2:
    print("\n--- GridSearchCV ROC Curves ---")
    plt.figure(figsize=(8, 6))

    fpr_grid_80, tpr_grid_80, _ = roc_curve(y_test_20_encoded, y_prob_80)
    roc_auc_grid_80 = auc(fpr_grid_80, tpr_grid_80)
    plt.plot(fpr_grid_80, tpr_grid_80, color='darkorange', lw=2, label='GridSearchCV 80/20 (area = %0.2f)' % roc_auc_grid_80)

```

```
fpr_grid_70, tpr_grid_70, _ = roc_curve(y_test_30_encoded, y_prob_70)
roc_auc_grid_70 = auc(fpr_grid_70, tpr_grid_70)
plt.plot(fpr_grid_70, tpr_grid_70, color='green', lw=2, label='GridSearchCV 70/30 (area = %0.2f)' % roc_auc_grid_70)

fpr_grid_60, tpr_grid_60, _ = roc_curve(y_test_40_encoded, y_prob_60)
roc_auc_grid_60 = auc(fpr_grid_60, tpr_grid_60)
plt.plot(fpr_grid_60, tpr_grid_60, color='red', lw=2, label='GridSearchCV 60/40 (area = %0.2f)' % roc_auc_grid_60)

plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve - GridSearchCV KNN')
plt.legend(loc="lower right")
plt.show()
```



```
Performing GridSearchCV on KNN (80/20 split)...
Best Parameters (80/20 split): {'knn__metric': 'manhattan', 'knn__n_neighbors': 7, 'knn__weights': 'distance'}
Test Accuracy (80/20 split): 0.8777777777777778
```

```
Performing GridSearchCV on KNN (70/30 split)...
Best Parameters (70/30 split): {'knn__metric': 'euclidean', 'knn__n_neighbors': 7, 'knn__weights': 'distance'}
Test Accuracy (70/30 split): 0.825925925925926
```


```
Performing GridSearchCV on KNN (60/40 split)...
Best Parameters (60/40 split): {'knn__metric': 'manhattan', 'knn__n_neighbors': 12, 'knn__weights': 'uniform'}
Test Accuracy (60/40 split): 0.8388888888888889
```

--- GridSearchCV Classification Reports and Confusion Matrices ---

```
GridSearchCV Classification Report (80/20 split KNN):
precision    recall    f1-score   support
      Besni       0.95      0.80      0.87      90
      Kecimen      0.83      0.96      0.89      90

      accuracy          0.88      180
      macro avg       0.89      0.88      0.88      180
      weighted avg    0.89      0.88      0.88      180
```


GridSearchCV Confusion Matrix (80/20 split KNN):



```
GridSearchCV Classification Report (70/30 split KNN):
precision    recall    f1-score   support
      Besni       0.87      0.76      0.81     135
      Kecimen      0.79      0.89      0.84     135

      accuracy          0.83      270
      macro avg       0.83      0.83      0.83     270
      weighted avg    0.83      0.83      0.83     270
```

GridSearchCV Confusion Matrix (70/30 split KNN):



```
print("Training Naive Bayes model...")
```

```
nb_pipeline = Pipeline(steps=[('classifier', GaussianNB())])

nb_pipeline.fit(X_train_80, y_train_80)

y_pred_nb = nb_pipeline.predict(X_test_20)

print("\nNaive Bayes Accuracy (80/20 split):", accuracy_score(y_test_20, y_pred_nb))

n_classes = len(np.unique(y_train_80))
```

```


if n_classes > 2:
    print("\nClassification Report (Naive Bayes):\n", classification_report(y_test_20, y_pred_nb))
    print("\nConfusion Matrix (Naive Bayes):\n", confusion_matrix(y_test_20, y_pred_nb))

else:
    print("\nClassification Report (Naive Bayes):\n", classification_report(y_test_20, y_pred_nb, zero_division=0))
    print("\nConfusion Matrix (Naive Bayes):\n", confusion_matrix(y_test_20, y_pred_nb))


if n_classes == 2:
    y_prob_nb = nb_pipeline.predict_proba(X_test_20)[:, 1]
    y_test_20_encoded_for_roc = y_test_20
    fpr, tpr, thresholds = roc_curve(y_test_20_encoded_for_roc, y_prob_nb)
    roc_auc = auc(fpr, tpr)

    plt.figure(figsize=(8, 6))
    plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver Operating Characteristic (ROC) Curve - Naive Bayes')
    plt.legend(loc="lower right")
    plt.show()

```



Confusion Matrix (Naive Bayes):
[[64 120] [4 86]] Receiver Operating Characteristic (ROC) Curve - GridSearchCV KNN
Receiver Operating Characteristic (ROC) Curve - Naive Bayes



Gaussian Naive Bayes Hyperparameter Justification:

- `[var_smoothing]` (Additive smoothing): Adds a small value to variances to improve numerical stability and handle zero frequencies.

```

from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, roc_curve, auc
import matplotlib.pyplot as plt

```

```
import seaborn as sns
from scipy.stats import loguniform

nb = GaussianNB()

print("\n--- Starting GridSearchCV for GaussianNB (80/20 split) ---")
param_grid_nb = {}
grid_search_nb_80 = GridSearchCV(estimator=nb, param_grid=param_grid_nb, cv=5, scoring='accuracy', n_jobs=-1)
grid_search_nb_80.fit(X_train_80, y_train_80)
print("\n--- GridSearchCV (80/20 split GaussianNB) Complete ---")
print(f"Best parameters found: {grid_search_nb_80.best_params_}")
print(f"Best cross-validation accuracy: {grid_search_nb_80.best_score_:.4f}")
y_pred_grid_nb_80 = grid_search_nb_80.predict(X_test_20)
print(f"\nGridSearchCV Test Accuracy: {accuracy_score(y_test_20, y_pred_grid_nb_80):.4f}")

print("\n--- Starting GridSearchCV for GaussianNB (70/30 split) ---")
grid_search_nb_70 = GridSearchCV(estimator=nb, param_grid=param_grid_nb, cv=5, scoring='accuracy', n_jobs=-1)
grid_search_nb_70.fit(X_train_70, y_train_70)
print("\n--- GridSearchCV (70/30 split GaussianNB) Complete ---")
print(f"Best parameters found: {grid_search_nb_70.best_params_}")
print(f"Best cross-validation accuracy: {grid_search_nb_70.best_score_:.4f}")
y_pred_grid_nb_70 = grid_search_nb_70.predict(X_test_30)
print(f"\nGridSearchCV Test Accuracy: {accuracy_score(y_test_30, y_pred_grid_nb_70):.4f}")

print("\n--- Starting GridSearchCV for GaussianNB (60/40 split) ---")
grid_search_nb_60 = GridSearchCV(estimator=nb, param_grid=param_grid_nb, cv=5, scoring='accuracy', n_jobs=-1)
grid_search_nb_60.fit(X_train_60, y_train_60)
print("\n--- GridSearchCV (60/40 split GaussianNB) Complete ---")
print(f"Best parameters found: {grid_search_nb_60.best_params_}")
print(f"Best cross-validation accuracy: {grid_search_nb_60.best_score_:.4f}")
y_pred_grid_nb_60 = grid_search_nb_60.predict(X_test_40)
print(f"\nGridSearchCV Test Accuracy: {accuracy_score(y_test_40, y_pred_grid_nb_60):.4f}")

print("\n--- Starting RandomizedSearchCV for GaussianNB (80/20 split) ---")
param_dist_nb = {
    'var_smoothing': loguniform(1e-9, 1e-2)
}
random_search_nb_80 = RandomizedSearchCV(
    estimator=nb,
    param_distributions=param_dist_nb,
    n_iter=10,
    cv=5,
    scoring='accuracy',
    random_state=42,
    n_jobs=-1
)
random_search_nb_80.fit(X_train_80, y_train_80)
print("\n--- RandomizedSearchCV (80/20 split GaussianNB) Complete ---")
print(f"Best parameters found: {random_search_nb_80.best_params_}")
print(f"Best cross-validation accuracy: {random_search_nb_80.best_score_:.4f}")
y_pred_random_nb_80 = random_search_nb_80.predict(X_test_20)
print(f"\nRandomizedSearchCV Test Accuracy: {accuracy_score(y_test_20, y_pred_random_nb_80):.4f}")

print("\n--- Starting RandomizedSearchCV for GaussianNB (70/30 split) ---")
random_search_nb_70 = RandomizedSearchCV(
    estimator=nb,
    param_distributions=param_dist_nb,
    n_iter=10,
    cv=5,
    scoring='accuracy',
    random_state=42,
    n_jobs=-1
)
random_search_nb_70.fit(X_train_70, y_train_70)
print("\n--- RandomizedSearchCV (70/30 split GaussianNB) Complete ---")
print(f"Best parameters found: {random_search_nb_70.best_params_}")
print(f"Best cross-validation accuracy: {random_search_nb_70.best_score_:.4f}")
y_pred_random_nb_70 = random_search_nb_70.predict(X_test_30)
print(f"\nRandomizedSearchCV Test Accuracy: {accuracy_score(y_test_30, y_pred_random_nb_70):.4f}")

print("\n--- Starting RandomizedSearchCV for GaussianNB (60/40 split) ---")
random_search_nb_60 = RandomizedSearchCV(
    estimator=nb,
    param_distributions=param_dist_nb,
    n_iter=10,
    cv=5,
    scoring='accuracy',
    random_state=42,
    n_jobs=-1
)
random_search_nb_60.fit(X_train_60, y_train_60)
```

```

print("\n--- RandomizedSearchCV (60/40 split GaussianNB) Complete ---")
print(f"Best parameters found: {random_search_nb_60.best_params_}")
print(f"Best cross-validation accuracy: {random_search_nb_60.best_score_:.4f}")
y_pred_random_nb_60 = random_search_nb_60.predict(X_test_40)
print(f"\nRandomizedSearchCV Test Accuracy: {accuracy_score(y_test_40, y_pred_random_nb_60):.4f}")


print("\nGridSearchCV Best Estimator Classification Report (80/20 split):")
print(classification_report(y_test_20, y_pred_grid_nb_80, target_names=label_encoder.classes_))
print("\nGridSearchCV Best Estimator Confusion Matrix (80/20 split):")
cm_grid_nb_80 = confusion_matrix(y_test_20, y_pred_grid_nb_80)
sns.heatmap(cm_grid_nb_80, annot=True, fmt='d', cmap='Blues', xticklabels=label_encoder.classes_, yticklabels=label_encoder.classes_)
plt.title("GridSearchCV GaussianNB - Confusion Matrix (80/20 Split)")
plt.ylabel('Actual'); plt.xlabel('Predicted')
plt.show()

print("\nGridSearchCV Best Estimator Classification Report (70/30 split):")
print(classification_report(y_test_30, y_pred_grid_nb_70, target_names=label_encoder.classes_))
print("\nGridSearchCV Best Estimator Confusion Matrix (70/30 split):")
cm_grid_nb_70 = confusion_matrix(y_test_30, y_pred_grid_nb_70)
sns.heatmap(cm_grid_nb_70, annot=True, fmt='d', cmap='Blues', xticklabels=label_encoder.classes_, yticklabels=label_encoder.classes_)
plt.title("GridSearchCV GaussianNB - Confusion Matrix (70/30 Split)")
plt.ylabel('Actual'); plt.xlabel('Predicted')
plt.show()

print("\nGridSearchCV Best Estimator Classification Report (60/40 split):")
print(classification_report(y_test_40, y_pred_grid_nb_60, target_names=label_encoder.classes_))
print("\nGridSearchCV Best Estimator Confusion Matrix (60/40 split):")
cm_grid_nb_60 = confusion_matrix(y_test_40, y_pred_grid_nb_60)
sns.heatmap(cm_grid_nb_60, annot=True, fmt='d', cmap='Blues', xticklabels=label_encoder.classes_, yticklabels=label_encoder.classes_)
plt.title("GridSearchCV GaussianNB - Confusion Matrix (60/40 Split)")
plt.ylabel('Actual'); plt.xlabel('Predicted')
plt.show()


print("\nRandomizedSearchCV Best Estimator Classification Report (80/20 split):")
print(classification_report(y_test_20, y_pred_random_nb_80, target_names=label_encoder.classes_))
print("\nRandomizedSearchCV Best Estimator Confusion Matrix (80/20 split):")
cm_random_nb_80 = confusion_matrix(y_test_20, y_pred_random_nb_80)
sns.heatmap(cm_random_nb_80, annot=True, fmt='d', cmap='Blues', xticklabels=label_encoder.classes_, yticklabels=label_encoder.classes_)
plt.title("RandomizedSearchCV GaussianNB - Confusion Matrix (80/20 Split)")
plt.ylabel('Actual'); plt.xlabel('Predicted')
plt.show()

print("\nRandomizedSearchCV Best Estimator Classification Report (70/30 split):")
print(classification_report(y_test_30, y_pred_random_nb_70, target_names=label_encoder.classes_))
print("\nRandomizedSearchCV Best Estimator Confusion Matrix (70/30 split):")
cm_random_nb_70 = confusion_matrix(y_test_30, y_pred_random_nb_70)
sns.heatmap(cm_random_nb_70, annot=True, fmt='d', cmap='Blues', xticklabels=label_encoder.classes_, yticklabels=label_encoder.classes_)
plt.title("RandomizedSearchCV GaussianNB - Confusion Matrix (70/30 Split)")
plt.ylabel('Actual'); plt.xlabel('Predicted')
plt.show()

print("\nRandomizedSearchCV Best Estimator Classification Report (60/40 split):")
print(classification_report(y_test_40, y_pred_random_nb_60, target_names=label_encoder.classes_))
print("\nRandomizedSearchCV Best Estimator Confusion Matrix (60/40 split):")
cm_random_nb_60 = confusion_matrix(y_test_40, y_pred_random_nb_60)
sns.heatmap(cm_random_nb_60, annot=True, fmt='d', cmap='Blues', xticklabels=label_encoder.classes_, yticklabels=label_encoder.classes_)
plt.title("RandomizedSearchCV GaussianNB - Confusion Matrix (60/40 Split)")
plt.ylabel('Actual'); plt.xlabel('Predicted')
plt.show()


if len(label_encoder.classes_) == 2:
    y_test_20_encoded = y_test_20
    y_test_30_encoded = y_test_30
    y_test_40_encoded = y_test_40

    print("\nGridSearchCV Best Estimator ROC Curve:")
    y_prob_grid_nb_80 = grid_search_nb_80.best_estimator_.predict_proba(X_test_20)[:, 1]
    fpr_grid_nb_80, tpr_grid_nb_80, thresholds_grid_nb_80 = roc_curve(y_test_20_encoded, y_prob_grid_nb_80)
    roc_auc_grid_nb_80 = auc(fpr_grid_nb_80, tpr_grid_nb_80)

    plt.figure(figsize=(8, 6))
    plt.plot(fpr_grid_nb_80, tpr_grid_nb_80, color='darkorange', lw=2, label='80/20 ROC curve (area = %0.2f)' % roc_auc_grid_nb_80)
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC Curve - GridSearchCV (GaussianNB) 80/20 Split')
    plt.legend(loc="lower right")

```

```
plt.show()

y_prob_grid_nb_70 = grid_search_nb_70.best_estimator_.predict_proba(X_test_30)[:, 1]
fpr_grid_nb_70, tpr_grid_nb_70, thresholds_grid_nb_70 = roc_curve(y_test_30_encoded, y_prob_grid_nb_70)
roc_auc_grid_nb_70 = auc(fpr_grid_nb_70, tpr_grid_nb_70)

plt.figure(figsize=(8, 6))
plt.plot(fpr_grid_nb_70, tpr_grid_nb_70, color='darkorange', lw=2, label='70/30 ROC curve (area = %0.2f)' % roc_auc_grid_nb_70)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - GridSearchCV (GaussianNB) 70/30 Split')
plt.legend(loc="lower right")
plt.show()

y_prob_grid_nb_60 = grid_search_nb_60.best_estimator_.predict_proba(X_test_40)[:, 1]
fpr_grid_nb_60, tpr_grid_nb_60, thresholds_grid_nb_60 = roc_curve(y_test_40_encoded, y_prob_grid_nb_60)
roc_auc_grid_nb_60 = auc(fpr_grid_nb_60, tpr_grid_nb_60)

plt.figure(figsize=(8, 6))
plt.plot(fpr_grid_nb_60, tpr_grid_nb_60, color='darkorange', lw=2, label='60/40 ROC curve (area = %0.2f)' % roc_auc_grid_nb_60)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - GridSearchCV (GaussianNB) 60/40 Split')
plt.legend(loc="lower right")
plt.show()

print("\nRandomizedSearchCV Best Estimator ROC Curve:")
y_prob_random_nb_80 = random_search_nb_80.best_estimator_.predict_proba(X_test_20)[:, 1]
fpr_random_nb_80, tpr_random_nb_80, thresholds_random_nb_80 = roc_curve(y_test_20_encoded, y_prob_random_nb_80)
roc_auc_random_nb_80 = auc(fpr_random_nb_80, tpr_random_nb_80)

plt.figure(figsize=(8, 6))
plt.plot(fpr_random_nb_80, tpr_random_nb_80, color='darkorange', lw=2, label='80/20 ROC curve (area = %0.2f)' % roc_auc_random_nb_80)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - RandomizedSearchCV (GaussianNB) 80/20 Split')
plt.legend(loc="lower right")
plt.show()

y_prob_random_nb_70 = random_search_nb_70.best_estimator_.predict_proba(X_test_30)[:, 1]
fpr_random_nb_70, tpr_random_nb_70, thresholds_random_nb_70 = roc_curve(y_test_30_encoded, y_prob_random_nb_70)
roc_auc_random_nb_70 = auc(fpr_random_nb_70, tpr_random_nb_70)

plt.figure(figsize=(8, 6))
plt.plot(fpr_random_nb_70, tpr_random_nb_70, color='darkorange', lw=2, label='70/30 ROC curve (area = %0.2f)' % roc_auc_random_nb_70)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - RandomizedSearchCV (GaussianNB) 70/30 Split')
plt.legend(loc="lower right")
plt.show()

y_prob_random_nb_60 = random_search_nb_60.best_estimator_.predict_proba(X_test_40)[:, 1]
fpr_random_nb_60, tpr_random_nb_60, thresholds_random_nb_60 = roc_curve(y_test_40_encoded, y_prob_random_nb_60)
roc_auc_random_nb_60 = auc(fpr_random_nb_60, tpr_random_nb_60)

plt.figure(figsize=(8, 6))
plt.plot(fpr_random_nb_60, tpr_random_nb_60, color='darkorange', lw=2, label='60/40 ROC curve (area = %0.2f)' % roc_auc_random_nb_60)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - RandomizedSearchCV (GaussianNB) 60/40 Split')
plt.legend(loc="lower right")
plt.show()
```



```

--- Starting GridSearchCV for GaussianNB (80/20 split) ---
--- GridSearchCV (80/20 split GaussianNB) Complete ---
Best parameters found: {}
Best cross-validation accuracy: 0.8361

GridSearchCV Test Accuracy: 0.8333

--- Starting GridSearchCV for GaussianNB (70/30 split) ---
--- GridSearchCV (70/30 split GaussianNB) Complete ---
Best parameters found: {}
Best cross-validation accuracy: 0.8540

GridSearchCV Test Accuracy: 0.7963

--- Starting GridSearchCV for GaussianNB (60/40 split) ---
--- GridSearchCV (60/40 split GaussianNB) Complete ---
Best parameters found: {}
Best cross-validation accuracy: 0.8537

GridSearchCV Test Accuracy: 0.8139

--- Starting RandomizedSearchCV for GaussianNB (80/20 split) ---
--- RandomizedSearchCV (80/20 split GaussianNB) Complete ---
Best parameters found: {'var_smoothing': np.float64(0.004518560951024112)}
Best cross-validation accuracy: 0.8375

RandomizedSearchCV Test Accuracy: 0.8333

--- Starting RandomizedSearchCV for GaussianNB (70/30 split) ---
--- RandomizedSearchCV (70/30 split GaussianNB) Complete ---
Best parameters found: {'var_smoothing': np.float64(0.004518560951024112)}
Best cross-validation accuracy: 0.8556

RandomizedSearchCV Test Accuracy: 0.8000

--- Starting RandomizedSearchCV for GaussianNB (60/40 split) ---
--- RandomizedSearchCV (60/40 split GaussianNB) Complete ---
Best parameters found: {'var_smoothing': np.float64(4.185822729546971e-07)}
Best cross-validation accuracy: 0.8537


RandomizedSearchCV Test Accuracy: 0.8139

```

GridSearchCV Best Estimator Classification Report (80/20 split):

	precision	recall	f1-score	support
Besni	0.94	0.71	0.81	90
Kecimen	0.77	0.96	0.85	90
accuracy			0.83	180
macro avg	0.85	0.83	0.83	180
weighted avg	0.85	0.83	0.83	180

GridSearchCV Best Estimator Confusion Matrix (80/20 split):



GridSearchCV Best Estimator Classification Report (70/30 split):